

MLP Architectures

In [5]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
import warnings
warnings.filterwarnings("ignore")
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [6]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [7]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [8]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [9]:

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [10]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [11]:

```
# An example data point
print(X_train[0])
```

[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	82	
	82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0															

In [12]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - Xmin)/(Xmax-Xmin) = X/255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [13]:

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [14]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [15]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
```

```

# Dense(10),
# Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [16]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [17]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```

packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version. Instructions for updating:
Colocations handled automatically by placer.

In [18]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format

# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From C:\Users\Shashank\Anaconda3\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 1s 24us/step - loss: 1.2839 - acc: 0.6860 - val_loss: 0.8109 - val_acc: 0.8381

Epoch 2/20

60000/60000 [=====] - 1s 20us/step - loss: 0.7170 - acc: 0.8403 - val_loss: 0.6064 - val_acc: 0.8646

Epoch 3/20

60000/60000 [=====] - 1s 20us/step - loss: 0.5871 - acc: 0.8588 - val_loss: 0.5241 - val_acc: 0.8770

Epoch 4/20

60000/60000 [=====] - 1s 20us/step - loss: 0.5253 - acc: 0.8682 - val_loss: 0.4787 - val_acc: 0.8830

Epoch 5/20

60000/60000 [=====] - 1s 21us/step - loss: 0.4876 - acc: 0.8754 - val_loss: 0.4490 - val_acc: 0.8878

Epoch 6/20

60000/60000 [=====] - 1s 20us/step - loss: 0.4618 - acc: 0.8801 - val_loss: 0.4275 - val_acc: 0.8922

Epoch 7/20

60000/60000 [=====] - 1s 19us/step - loss: 0.4426 - acc: 0.8835 -

```

00000/00000 [=====] - 1s 19us/step - loss: 0.4121 - acc: 0.8945
val_loss: 0.4121 - val_acc: 0.8945
Epoch 8/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4277 - acc: 0.8865 -
val_loss: 0.3990 - val_acc: 0.8968
Epoch 9/20
60000/60000 [=====] - 1s 20us/step - loss: 0.4157 - acc: 0.8892 -
val_loss: 0.3886 - val_acc: 0.8980
Epoch 10/20
60000/60000 [=====] - 1s 21us/step - loss: 0.4057 - acc: 0.8916 -
val_loss: 0.3800 - val_acc: 0.8988
Epoch 11/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3973 - acc: 0.8935 -
val_loss: 0.3727 - val_acc: 0.9011
Epoch 12/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3900 - acc: 0.8949 -
val_loss: 0.3663 - val_acc: 0.9026
Epoch 13/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3837 - acc: 0.8964 -
val_loss: 0.3606 - val_acc: 0.9036
Epoch 14/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3780 - acc: 0.8973 -
val_loss: 0.3561 - val_acc: 0.9055
Epoch 15/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3730 - acc: 0.8986 -
val_loss: 0.3517 - val_acc: 0.9071
Epoch 16/20
60000/60000 [=====] - 1s 22us/step - loss: 0.3685 - acc: 0.8996 -
val_loss: 0.3475 - val_acc: 0.9078
Epoch 17/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3644 - acc: 0.9005 -
val_loss: 0.3439 - val_acc: 0.9086
Epoch 18/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3606 - acc: 0.9014 -
val_loss: 0.3409 - val_acc: 0.9087
Epoch 19/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3572 - acc: 0.9019 -
val_loss: 0.3379 - val_acc: 0.9098
Epoch 20/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3541 - acc: 0.9025 -
val_loss: 0.3351 - val_acc: 0.9104

```

In [19]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.33509153289198873

Test accuracy: 0.9104



MLP + Sigmoid activation + SGDOptimizer

In [21]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [22]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 126us/step - loss: 2.2614 - acc: 0.2649 -
val_loss: 2.2083 - val_acc: 0.3769
Epoch 2/20
60000/60000 [=====] - 8s 133us/step - loss: 2.1583 - acc: 0.4630 -
val_loss: 2.0966 - val_acc: 0.4593
Epoch 3/20
60000/60000 [=====] - 8s 138us/step - loss: 2.0264 - acc: 0.5711 -
val_loss: 1.9356 - val_acc: 0.5871
```



```

Epoch 4/20
60000/60000 [=====] - 8s 133us/step - loss: 1.8409 - acc: 0.6328 -
val_loss: 1.7199 - val_acc: 0.6586
Epoch 5/20
60000/60000 [=====] - 8s 141us/step - loss: 1.6112 - acc: 0.6831 -
val_loss: 1.4777 - val_acc: 0.7173
Epoch 6/20
60000/60000 [=====] - 8s 134us/step - loss: 1.3782 - acc: 0.7293 -
val_loss: 1.2556 - val_acc: 0.7624
Epoch 7/20
60000/60000 [=====] - 8s 130us/step - loss: 1.1786 - acc: 0.7624 -
val_loss: 1.0760 - val_acc: 0.7861
Epoch 8/20
60000/60000 [=====] - 8s 130us/step - loss: 1.0225 - acc: 0.7862 -
val_loss: 0.9405 - val_acc: 0.8074
Epoch 9/20
60000/60000 [=====] - 9s 146us/step - loss: 0.9045 - acc: 0.8033 -
val_loss: 0.8384 - val_acc: 0.8173
Epoch 10/20
60000/60000 [=====] - 9s 156us/step - loss: 0.8146 - acc: 0.8161 -
val_loss: 0.7597 - val_acc: 0.8311
Epoch 11/20
60000/60000 [=====] - 9s 150us/step - loss: 0.7445 - acc: 0.8271 -
val_loss: 0.6976 - val_acc: 0.8395
Epoch 12/20
60000/60000 [=====] - 9s 151us/step - loss: 0.6888 - acc: 0.8359 -
val_loss: 0.6480 - val_acc: 0.8451
Epoch 13/20
60000/60000 [=====] - 9s 156us/step - loss: 0.6436 - acc: 0.8435 -
val_loss: 0.6074 - val_acc: 0.8512
Epoch 14/20
60000/60000 [=====] - 10s 165us/step - loss: 0.6063 - acc: 0.8501 - val_l
oss: 0.5738 - val_acc: 0.8578
Epoch 15/20
60000/60000 [=====] - 9s 153us/step - loss: 0.5753 - acc: 0.8561 -
val_loss: 0.5447 - val_acc: 0.8630
Epoch 16/20
60000/60000 [=====] - 8s 137us/step - loss: 0.5490 - acc: 0.8601 -
val_loss: 0.5213 - val_acc: 0.8671
Epoch 17/20
60000/60000 [=====] - 10s 162us/step - loss: 0.5265 - acc: 0.8646 - val_l
oss: 0.5004 - val_acc: 0.8711
Epoch 18/20
60000/60000 [=====] - 9s 158us/step - loss: 0.5073 - acc: 0.8681 -
val_loss: 0.4827 - val_acc: 0.8742
Epoch 19/20
60000/60000 [=====] - 8s 128us/step - loss: 0.4904 - acc: 0.8706 -
val_loss: 0.4661 - val_acc: 0.8774
Epoch 20/20
60000/60000 [=====] - 8s 135us/step - loss: 0.4757 - acc: 0.8739 -
val_loss: 0.4528 - val_acc: 0.8797

```

In [23]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.4527508878707886
Test accuracy: 0.8797



In [24]:


```
w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + Sigmoid activation + ADAM

In [25]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 202us/step - loss: 0.5493 - acc: 0.8559 - val_loss: 0.2645 - val_acc: 0.9198
Epoch 2/20
60000/60000 [=====] - 11s 180us/step - loss: 0.2243 - acc: 0.9337 - val_loss: 0.1883 - val_acc: 0.9448
Epoch 3/20
60000/60000 [=====] - 10s 173us/step - loss: 0.1640 - acc: 0.9516 - val_loss: 0.1444 - val_acc: 0.9579
Epoch 4/20
60000/60000 [=====] - 12s 192us/step - loss: 0.1267 - acc: 0.9626 - val_loss: 0.1229 - val_acc: 0.9623
Epoch 5/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0980 - acc: 0.9711 - val_loss: 0.1009 - val_acc: 0.9676
Epoch 6/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0788 - acc: 0.9763 - val_loss: 0.0900 - val_acc: 0.9729
Epoch 7/20
60000/60000 [=====] - 12s 199us/step - loss: 0.0635 - acc: 0.9811 - val_loss: 0.0803 - val_acc: 0.9750
Epoch 8/20
60000/60000 [=====] - 12s 207us/step - loss: 0.0514 - acc: 0.9849 - val_loss: 0.0693 - val_acc: 0.9783
```

```

Epoch 9/20
60000/60000 [=====] - 11s 192us/step - loss: 0.0415 - acc: 0.9881 - val_loss: 0.0739 - val_acc: 0.9775
Epoch 10/20
60000/60000 [=====] - 12s 193us/step - loss: 0.0333 - acc: 0.9905 - val_loss: 0.0722 - val_acc: 0.9774
Epoch 11/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0267 - acc: 0.9931 - val_loss: 0.0627 - val_acc: 0.9805
Epoch 12/20
60000/60000 [=====] - 12s 207us/step - loss: 0.0223 - acc: 0.9942 - val_loss: 0.0670 - val_acc: 0.9784
Epoch 13/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0174 - acc: 0.9957 - val_loss: 0.0674 - val_acc: 0.9798
Epoch 14/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0132 - acc: 0.9972 - val_loss: 0.0681 - val_acc: 0.9790 - loss
Epoch 15/20
60000/60000 [=====] - 12s 200us/step - loss: 0.0113 - acc: 0.9973 - val_loss: 0.0680 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0089 - acc: 0.9981 - val_loss: 0.0673 - val_acc: 0.9815
Epoch 17/20
60000/60000 [=====] - 14s 226us/step - loss: 0.0078 - acc: 0.9981 - val_loss: 0.0689 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] - 14s 225us/step - loss: 0.0051 - acc: 0.9991 - val_loss: 0.0676 - val_acc: 0.9809
Epoch 19/20
60000/60000 [=====] - 15s 245us/step - loss: 0.0045 - acc: 0.9993 - val_loss: 0.0680 - val_acc: 0.9815acc:
Epoch 20/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0036 - acc: 0.9992 - val_loss: 0.0669 - val_acc: 0.9826s - loss: 0.0037 - a

```

In [26]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06692838039992276

Test accuracy: 0.9826





In [27]:


```
w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + ReLU +SGD

In [28]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [29]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 159us/step - loss: 0.7577 - acc: 0.7888 - val_loss: 0.3863 - val_acc: 0.8958
Epoch 2/20
60000/60000 [=====] - 8s 141us/step - loss: 0.3511 - acc: 0.9027 - val_loss: 0.2988 - val_acc: 0.9164
Epoch 3/20
60000/60000 [=====] - 8s 135us/step - loss: 0.2885 - acc: 0.9201 - val_loss: 0.2609 - val_acc: 0.9256
Epoch 4/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2545 - acc: 0.9292 - val_loss: 0.2368 - val_acc: 0.9329
Epoch 5/20
60000/60000 [=====] - 8s 138us/step - loss: 0.2310 - acc: 0.9362 - val_loss: 0.2190 - val_acc: 0.9367
Epoch 6/20
60000/60000 [=====] - 9s 153us/step - loss: 0.2130 - acc: 0.9410 - val_loss: 0.2041 - val_acc: 0.9406
Epoch 7/20
60000/60000 [=====] - 10s 160us/step - loss: 0.1984 - acc: 0.9447 - val_loss: 0.1943 - val_acc: 0.9429
Epoch 8/20
```

```

Epoch 8/20
60000/60000 [=====] - 8s 125us/step - loss: 0.1862 - acc: 0.9479 -
val_loss: 0.1833 - val_acc: 0.9454
Epoch 9/20
60000/60000 [=====] - 7s 123us/step - loss: 0.1757 - acc: 0.9510 -
val_loss: 0.1755 - val_acc: 0.9470
Epoch 10/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1664 - acc: 0.9538 -
val_loss: 0.1676 - val_acc: 0.9504
Epoch 11/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1580 - acc: 0.9561 -
val_loss: 0.1621 - val_acc: 0.9521
Epoch 12/20
60000/60000 [=====] - 8s 129us/step - loss: 0.1507 - acc: 0.9582 -
val_loss: 0.1564 - val_acc: 0.9543
Epoch 13/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1440 - acc: 0.9599 -
val_loss: 0.1533 - val_acc: 0.9551
Epoch 14/20
60000/60000 [=====] - 9s 148us/step - loss: 0.1379 - acc: 0.9618 -
val_loss: 0.1458 - val_acc: 0.9577
Epoch 15/20
60000/60000 [=====] - 8s 127us/step - loss: 0.1323 - acc: 0.9632 -
val_loss: 0.1409 - val_acc: 0.9588
Epoch 16/20
60000/60000 [=====] - 7s 124us/step - loss: 0.1270 - acc: 0.9650 -
val_loss: 0.1382 - val_acc: 0.9603
Epoch 17/20
60000/60000 [=====] - 8s 132us/step - loss: 0.1224 - acc: 0.9661 -
val_loss: 0.1326 - val_acc: 0.9602
Epoch 18/20
60000/60000 [=====] - 8s 134us/step - loss: 0.1177 - acc: 0.9676 -
val_loss: 0.1302 - val_acc: 0.9620
Epoch 19/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1135 - acc: 0.9686 -
val_loss: 0.1268 - val_acc: 0.9634
Epoch 20/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1098 - acc: 0.9695 -
val_loss: 0.1243 - val_acc: 0.9635

```

In [30]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12430026308484375
Test accuracy: 0.9635



In [31]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + ReLU + ADAM

In [32]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
dense_15 (Dense)	(None, 128)	65664
dense_16 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 172us/step - loss: 0.2312 - acc: 0.9304 - val_loss: 0.1250 - val_acc: 0.9617
Epoch 2/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0854 - acc: 0.9745 - val_loss: 0.0906 - val_acc: 0.9711
Epoch 3/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0537 - acc: 0.9836 - val_loss: 0.0872 - val_acc: 0.9728
Epoch 4/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0362 - acc: 0.9889 - val_loss: 0.0671 - val_acc: 0.9790
Epoch 5/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0262 - acc: 0.9912 - val_loss: 0.0690 - val_acc: 0.9775
Epoch 6/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0202 - acc: 0.9939 - val_loss: 0.0705 - val_acc: 0.9782
Epoch 7/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0153 - acc: 0.9948 - val_loss: 0.0863 - val_acc: 0.9764
Epoch 8/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0145 - acc: 0.9951 - val_loss: 0.0897 - val_acc: 0.9767
Epoch 9/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0148 - acc: 0.9948 - val_loss: 0.0759 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 10s 171us/step - loss: 0.0115 - acc: 0.9962 - val_loss: 0.0921 - val_acc: 0.9767
Epoch 11/20
60000/60000 [=====] - 10s 167us/step - loss: 0.0093 - acc: 0.9970 - val_loss: 0.1128 - val_acc: 0.9744
```

```

Epoch 12/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0123 - acc: 0.9957 - val_loss: 0.0815 - val_acc: 0.9796
Epoch 13/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0105 - acc: 0.9967 - val_loss: 0.0806 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0077 - acc: 0.9976 - val_loss: 0.0940 - val_acc: 0.9785
Epoch 15/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0116 - acc: 0.9962 - val_loss: 0.0885 - val_acc: 0.9801
Epoch 16/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0086 - acc: 0.9973 - val_loss: 0.0961 - val_acc: 0.9789
Epoch 17/20
60000/60000 [=====] - 11s 184us/step - loss: 0.0069 - acc: 0.9977 - val_loss: 0.0918 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] - 10s 170us/step - loss: 0.0083 - acc: 0.9972 - val_loss: 0.0847 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 11s 186us/step - loss: 0.0052 - acc: 0.9983 - val_loss: 0.0894 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 11s 177us/step - loss: 0.0048 - acc: 0.9984 - val_loss: 0.1014 - val_acc: 0.9780

```

In [33]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10140154234402549

Test accuracy: 0.978





In [34]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [35]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_18 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_19 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [36]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 228us/step - loss: 0.3058 - acc: 0.9084 - val_loss: 0.2115 - val_acc: 0.9404
Epoch 2/20
60000/60000 [=====] - 12s 194us/step - loss: 0.1744 - acc: 0.9489 - val_loss: 0.1689 - val_acc: 0.9510
Epoch 3/20
60000/60000 [=====] - 12s 203us/step - loss: 0.1357 - acc: 0.9607 - val_loss: 0.1452 - val_acc: 0.9575
Epoch 4/20
60000/60000 [=====] - 11s 185us/step - loss: 0.1138 - acc: 0.9664 - val_loss: 0.1304 - val_acc: 0.9623
Epoch 5/20
60000/60000 [=====] - 12s 193us/step - loss: 0.0966 - acc: 0.9715 - val_loss: 0.1337 - val_acc: 0.9596
Epoch 6/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0808 - acc: 0.9755 - val_loss: 0.1103 - val_acc: 0.9667
```

```

Epoch 7/20
60000/60000 [=====] - 11s 191us/step - loss: 0.0691 - acc: 0.9788 - val_loss: 0.1132 - val_acc: 0.9652
Epoch 8/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0613 - acc: 0.9809 - val_loss: 0.1104 - val_acc: 0.9659
Epoch 9/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0529 - acc: 0.9832 - val_loss: 0.1111 - val_acc: 0.9670
Epoch 10/20
60000/60000 [=====] - 11s 185us/step - loss: 0.0457 - acc: 0.9860 - val_loss: 0.0973 - val_acc: 0.9706
Epoch 11/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0393 - acc: 0.9877 - val_loss: 0.1003 - val_acc: 0.9710
Epoch 12/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0348 - acc: 0.9891 - val_loss: 0.0970 - val_acc: 0.9718
Epoch 13/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0306 - acc: 0.9907 - val_loss: 0.0995 - val_acc: 0.9707
Epoch 14/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0265 - acc: 0.9919 - val_loss: 0.0940 - val_acc: 0.9741
Epoch 15/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0256 - acc: 0.9921 - val_loss: 0.1006 - val_acc: 0.9734
Epoch 16/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0235 - acc: 0.9925 - val_loss: 0.0918 - val_acc: 0.9727
Epoch 17/20
60000/60000 [=====] - 12s 203us/step - loss: 0.0208 - acc: 0.9934 - val_loss: 0.0930 - val_acc: 0.9752
Epoch 18/20
60000/60000 [=====] - 13s 212us/step - loss: 0.0205 - acc: 0.9930 - val_loss: 0.1020 - val_acc: 0.9724
Epoch 19/20
60000/60000 [=====] - 12s 198us/step - loss: 0.0194 - acc: 0.9937 - val_loss: 0.0953 - val_acc: 0.9744
Epoch 20/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0153 - acc: 0.9952 - val_loss: 0.1014 - val_acc: 0.9741

```

In [37]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10143673473168165
Test accuracy: 0.9741



In [38]:

```
w_after = model_batch.get_weights()

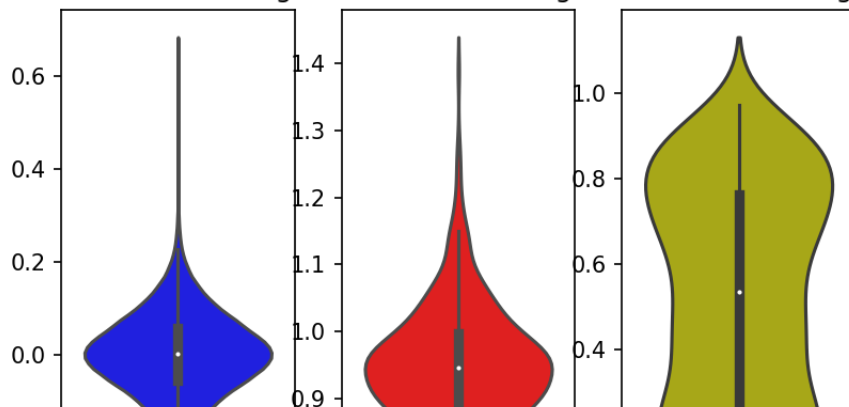
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model WeightsTrained model WeightsTrained model Weights





5. MLP + Dropout + AdamOptimizer

In [39]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.05, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:From C:\Users\Shashank\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_21 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [40]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 225us/step - loss: 0.6715 - acc: 0.7936 - val_loss: 0.2873 - val_acc: 0.9132s - loss
Epoch 2/20
60000/60000 [=====] - 12s 201us/step - loss: 0.4331 - acc: 0.8683 - val_loss: 0.2592 - val_acc: 0.9231
Epoch 3/20
60000/60000 [=====] - 12s 192us/step - loss: 0.3868 - acc: 0.8827 - val_loss: 0.2369 - val_acc: 0.9335
Epoch 4/20
60000/60000 [=====] - 11s 187us/step - loss: 0.3565 - acc: 0.8934 - val_loss: 0.2231 - val_acc: 0.9332
Epoch 5/20
60000/60000 [=====] - 11s 188us/step - loss: 0.3371 - acc: 0.8990 - val_loss: 0.2093 - val_acc: 0.9389
Epoch 6/20
60000/60000 [=====] - 12s 202us/step - loss: 0.3225 - acc: 0.9026 - val_loss: 0.2011 - val_acc: 0.9393
Epoch 7/20
60000/60000 [=====] - 13s 213us/step - loss: 0.3057 - acc: 0.9071 - val_loss: 0.1934 - val_acc: 0.9413
Epoch 8/20
60000/60000 [=====] - 12s 204us/step - loss: 0.2894 - acc: 0.9123 - val_loss: 0.1850 - val_acc: 0.9458
Epoch 9/20
60000/60000 [=====] - 12s 203us/step - loss: 0.2799 - acc: 0.9168 - val_loss: 0.1821 - val_acc: 0.9461
Epoch 10/20
60000/60000 [=====] - 12s 200us/step - loss: 0.2668 - acc: 0.9202 - val_loss: 0.1693 - val_acc: 0.9504
Epoch 11/20
60000/60000 [=====] - 12s 194us/step - loss: 0.2591 - acc: 0.9225 - val_loss: 0.1587 - val_acc: 0.9528
Epoch 12/20
60000/60000 [=====] - 12s 208us/step - loss: 0.2514 - acc: 0.9244 - val_loss: 0.1504 - val_acc: 0.9562
Epoch 13/20
60000/60000 [=====] - 13s 220us/step - loss: 0.2385 - acc: 0.9279 - val_loss: 0.1435 - val_acc: 0.9564
Epoch 14/20
60000/60000 [=====] - 13s 212us/step - loss: 0.2250 - acc: 0.9325 - val_loss: 0.1360 - val_acc: 0.9602
Epoch 15/20
60000/60000 [=====] - 12s 197us/step - loss: 0.2190 - acc: 0.9340 - val_loss: 0.1334 - val_acc: 0.9591
Epoch 16/20
60000/60000 [=====] - 14s 235us/step - loss: 0.2092 - acc: 0.9377 - val_loss: 0.1304 - val_acc: 0.9601
Epoch 17/20
60000/60000 [=====] - 12s 196us/step - loss: 0.1985 - acc: 0.9407 - val_loss: 0.1244 - val_acc: 0.9651
Epoch 18/20
60000/60000 [=====] - 12s 195us/step - loss: 0.1911 - acc: 0.9428 - val_loss: 0.1225 - val_acc: 0.9635
Epoch 19/20
60000/60000 [=====] - 13s 224us/step - loss: 0.1870 - acc: 0.9440 - val_loss: 0.1140 - val_acc: 0.9672
Epoch 20/20
60000/60000 [=====] - 13s 222us/step - loss: 0.1747 - acc: 0.9468 - val_loss: 0.1047 - val_acc: 0.9699

```

In [41]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```



```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10473509188406169

Test accuracy: 0.9699



In [42]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



(MLP + Dropout + AdamOptimizer(with different size of hidden layers(630,410)))

In [43]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(630, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(410, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 630)	494550
batch_normalization_5 (Batch Normalization)	(None, 630)	2520
dropout_3 (Dropout)	(None, 630)	0
dense_24 (Dense)	(None, 410)	258710
batch_normalization_6 (Batch Normalization)	(None, 410)	1640
dropout_4 (Dropout)	(None, 410)	0
dense_25 (Dense)	(None, 10)	4110

Total params: 761,530
Trainable params: 759,450
Non-trainable params: 2,080

In [44]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 24s 398us/step - loss: 0.6136 - acc: 0.8161 - val_loss: 0.2725 - val_acc: 0.9234
Epoch 2/20
60000/60000 [=====] - 22s 360us/step - loss: 0.3867 - acc: 0.8824 - val_loss: 0.2462 - val_acc: 0.9301
Epoch 3/20
60000/60000 [=====] - 20s 327us/step - loss: 0.3418 - acc: 0.8960 - val_loss: 0.2252 - val_acc: 0.9343
Epoch 4/20
60000/60000 [=====] - 23s 379us/step - loss: 0.3135 - acc: 0.9062 - val_loss: 0.2054 - val_acc: 0.9393
Epoch 5/20
60000/60000 [=====] - 21s 346us/step - loss: 0.2929 - acc: 0.9098 - val_loss: 0.1938 - val_acc: 0.9432
Epoch 6/20
60000/60000 [=====] - 20s 336us/step - loss: 0.2798 - acc: 0.9151 - val_loss: 0.1860 - val_acc: 0.9465
Epoch 7/20
60000/60000 [=====] - 18s 302us/step - loss: 0.2621 - acc: 0.9205 - val_loss: 0.1762 - val_acc: 0.9483
Epoch 8/20
60000/60000 [=====] - 18s 292us/step - loss: 0.2502 - acc: 0.9240 - val_loss: 0.1683 - val_acc: 0.9501
Epoch 9/20
60000/60000 [=====] - 18s 300us/step - loss: 0.2321 - acc: 0.9290 - val_loss: 0.1558 - val_acc: 0.9555
Epoch 10/20
60000/60000 [=====] - 18s 294us/step - loss: 0.2301 - acc: 0.9300 - val_loss: 0.1477 - val_acc: 0.9567
Epoch 11/20
60000/60000 [=====] - 19s 318us/step - loss: 0.2153 - acc: 0.9355 - val_loss: 0.1464 - val_acc: 0.9582
Epoch 12/20
60000/60000 [=====] - 24s 393us/step - loss: 0.2020 - acc: 0.9386 - val_loss: 0.1320 - val_acc: 0.9622
Epoch 13/20
60000/60000 [=====] - 19s 318us/step - loss: 0.1957 - acc: 0.9403 - val_loss: 0.1322 - val_acc: 0.9629
Epoch 14/20
60000/60000 [=====] - 19s 310us/step - loss: 0.1845 - acc: 0.9424 - val_loss: 0.1212 - val_acc: 0.9642
Epoch 15/20
60000/60000 [=====] - 19s 313us/step - loss: 0.1763 - acc: 0.9451 - val_loss: 0.1153 - val_acc: 0.9674
Epoch 16/20
60000/60000 [=====] - 19s 314us/step - loss: 0.1696 - acc: 0.9470 - val_loss: 0.1114 - val_acc: 0.9680
Epoch 17/20
60000/60000 [=====] - 26s 429us/step - loss: 0.1599 - acc: 0.9508 - val_loss: 0.1079 - val_acc: 0.9680
Epoch 18/20
60000/60000 [=====] - 26s 426us/step - loss: 0.1566 - acc: 0.9528 - val_loss: 0.0996 - val_acc: 0.9713
Epoch 19/20
60000/60000 [=====] - 29s 487us/step - loss: 0.1476 - acc: 0.9546 - val_loss: 0.0965 - val_acc: 0.9732 - acc: 0. - ETA: 0s - loss: 0.1472 - acc: 0.954
Epoch 20/20
60000/60000 [=====] - 27s 456us/step - loss: 0.1436 - acc: 0.9552 - val_loss: 0.0951 - val_acc: 0.9722
```

In [45]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09514947874806821
Test accuracy: 0.9722



In [46]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
```


```

ax = sns.violinplot(y=nl_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Dropout + AdamOptimizer(with different size of hidden layers(700,300))

In [47]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(600, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(200, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
=====		

dense_26 (Dense)	(None, 600)	471000
batch_normalization_7 (Batch Normalization)	(None, 600)	2400
dropout_5 (Dropout)	(None, 600)	0
dense_27 (Dense)	(None, 200)	120200
batch_normalization_8 (Batch Normalization)	(None, 200)	800
dropout_6 (Dropout)	(None, 200)	0
dense_28 (Dense)	(None, 10)	2010
=====		
Total params: 596,410		
Trainable params: 594,810		
Non-trainable params: 1,600		

In [48]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 22s 374us/step - loss: 0.6463 - acc: 0.8023 - val_loss: 0.2833 - val_acc: 0.9171
Epoch 2/20
60000/60000 [=====] - 20s 333us/step - loss: 0.4101 - acc: 0.8751 - val_loss: 0.2499 - val_acc: 0.9264
Epoch 3/20
60000/60000 [=====] - 21s 342us/step - loss: 0.3619 - acc: 0.8892 - val_loss: 0.2246 - val_acc: 0.9344
Epoch 4/20
60000/60000 [=====] - 19s 316us/step - loss: 0.3398 - acc: 0.8968 - val_loss: 0.2147 - val_acc: 0.9355
Epoch 5/20
60000/60000 [=====] - 17s 283us/step - loss: 0.3143 - acc: 0.9050 - val_loss: 0.1972 - val_acc: 0.9407
Epoch 6/20
60000/60000 [=====] - 20s 327us/step - loss: 0.2991 - acc: 0.9093 - val_loss: 0.1898 - val_acc: 0.9421
Epoch 7/20
60000/60000 [=====] - 19s 309us/step - loss: 0.2867 - acc: 0.9137 - val_loss: 0.1823 - val_acc: 0.9467
Epoch 8/20
60000/60000 [=====] - 21s 357us/step - loss: 0.2698 - acc: 0.9180 - val_loss: 0.1729 - val_acc: 0.9493
Epoch 9/20
60000/60000 [=====] - 21s 354us/step - loss: 0.2630 - acc: 0.9201 - val_loss: 0.1623 - val_acc: 0.9530
Epoch 10/20
60000/60000 [=====] - 20s 334us/step - loss: 0.2475 - acc: 0.9257 - val_loss: 0.1595 - val_acc: 0.9512
Epoch 11/20
60000/60000 [=====] - 19s 315us/step - loss: 0.2392 - acc: 0.9279 - val_loss: 0.1493 - val_acc: 0.9560
Epoch 12/20
60000/60000 [=====] - 19s 316us/step - loss: 0.2280 - acc: 0.9298 - val_loss: 0.1408 - val_acc: 0.9582
Epoch 13/20
60000/60000 [=====] - 20s 338us/step - loss: 0.2202 - acc: 0.9327 - val_loss: 0.1369 - val_acc: 0.9593
Epoch 14/20
60000/60000 [=====] - 18s 307us/step - loss: 0.2103 - acc: 0.9371 - val_loss: 0.1296 - val_acc: 0.9606
Epoch 15/20
60000/60000 [=====] - 16s 268us/step - loss: 0.2016 - acc: 0.9384 - val_loss: 0.1240 - val_acc: 0.9620
Epoch 16/20
60000/60000 [=====] - 14s 240us/step - loss: 0.1926 - acc: 0.9409 - val_loss: 0.1196 - val_acc: 0.9642
Epoch 17/20
```

```
Epoch 17/20
60000/60000 [=====] - 14s 235us/step - loss: 0.1820 - acc: 0.9450 - val_loss: 0.1161 - val_acc: 0.9650
Epoch 18/20
60000/60000 [=====] - 15s 257us/step - loss: 0.1750 - acc: 0.9460 - val_loss: 0.1073 - val_acc: 0.9673
Epoch 19/20
60000/60000 [=====] - 16s 272us/step - loss: 0.1648 - acc: 0.9494 - val_loss: 0.1080 - val_acc: 0.9675
Epoch 20/20
60000/60000 [=====] - 19s 317us/step - loss: 0.1601 - acc: 0.9510 - val_loss: 0.1034 - val_acc: 0.9700
```

In [49]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10339480265732855

Test accuracy: 0.97



In [50]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout + Adam Optimizer with 3 hidden layers(layer1 :800,layer2 :430,layer3:160)

In [51]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(800, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```



```

model_drop.add(Dense(430, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(160, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 800)	628000
batch_normalization_9 (Batch Normalization)	(None, 800)	3200
dropout_7 (Dropout)	(None, 800)	0
dense_30 (Dense)	(None, 430)	344430
batch_normalization_10 (Batch Normalization)	(None, 430)	1720
dropout_8 (Dropout)	(None, 430)	0
dense_31 (Dense)	(None, 160)	68960
batch_normalization_11 (Batch Normalization)	(None, 160)	640
dropout_9 (Dropout)	(None, 160)	0
dense_32 (Dense)	(None, 10)	1610
Total params: 1,048,560		
Trainable params: 1,045,780		
Non-trainable params: 2,780		

In [52]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
nb_epoch=15
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/15
60000/60000 [=====] - 34s 570us/step - loss: 1.1707 - acc: 0.6244 - val_loss: 0.4072 - val_acc: 0.8815
Epoch 2/15
60000/60000 [=====] - 27s 454us/step - loss: 0.6552 - acc: 0.7899 - val_loss: 0.3481 - val_acc: 0.9004
Epoch 3/15
60000/60000 [=====] - 31s 524us/step - loss: 0.5489 - acc: 0.8288 - val_loss: 0.3212 - val_acc: 0.9074
Epoch 4/15
60000/60000 [=====] - 35s 591us/step - loss: 0.4965 - acc: 0.8451 - val_loss: 0.2922 - val_acc: 0.9145
Epoch 5/15
60000/60000 [=====] - 33s 542us/step - loss: 0.4678 - acc: 0.8557 - val_loss: 0.2880 - val_acc: 0.9176
Epoch 6/15
60000/60000 [=====] - 33s 558us/step - loss: 0.4391 - acc: 0.8657 - val_loss: 0.2738 - val_acc: 0.9211
Epoch 7/15
60000/60000 [=====] - 33s 558us/step - loss: 0.4131 - acc: 0.8740 - val_loss: 0.2632 - val_acc: 0.9238
Epoch 8/15
60000/60000 [=====] - 30s 495us/step - loss: 0.3994 - acc: 0.8792 - val_loss: 0.2498 - val_acc: 0.9286

```

```

oss: 0.2193 - val_acc: 0.9289
Epoch 9/15
60000/60000 [=====] - 29s 488us/step - loss: 0.3847 - acc: 0.8837 - val_loss: 0.2488 - val_acc: 0.9279
Epoch 10/15
60000/60000 [=====] - 28s 465us/step - loss: 0.3734 - acc: 0.8869 - val_loss: 0.2317 - val_acc: 0.9341
Epoch 11/15
60000/60000 [=====] - 32s 536us/step - loss: 0.3559 - acc: 0.8931 - val_loss: 0.2223 - val_acc: 0.9376
Epoch 12/15
60000/60000 [=====] - 38s 637us/step - loss: 0.3460 - acc: 0.8960 - val_loss: 0.2127 - val_acc: 0.9396
Epoch 13/15
60000/60000 [=====] - 30s 504us/step - loss: 0.3339 - acc: 0.9001 - val_loss: 0.2052 - val_acc: 0.9404 - a - ETA: 2
Epoch 14/15
60000/60000 [=====] - 26s 435us/step - loss: 0.3239 - acc: 0.9036 - val_loss: 0.2037 - val_acc: 0.9433
Epoch 15/15
60000/60000 [=====] - 27s 455us/step - loss: 0.3104 - acc: 0.9085 - val_loss: 0.1936 - val_acc: 0.9446

```

In [53]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.19363341980278492

Test accuracy: 0.9446





In [54]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [60]:

```
# http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

#If you get a ModuleNotFoundError error , install prettytable using: pip3 install prettytable

x = PrettyTable()
```

```

x.field_names = ["Model", "Accuracy"]
x.add_row(["Softmax classifier", 90.02])
x.add_row(["MLP+Sigmod activation+SGD Optimizer", 88.05])
x.add_row(["MLP+Sigmoid activation+ADAM", 97.95])
x.add_row(["MLP+ReLU+SGD", 98.46])
x.add_row(["MLP+ReLU+ADAM", 98.43])
x.add_row(["MLP+Batch Norm on hidden Layers+Adam Optimizer", 97.37])
x.add_row(["MLP+Dropout+AdamOptimizer", 96.92])
x.add_row(["MLP +Dropout + AdamOptimizer(different size of hidden layers(540,360))", 97.05])
x.add_row(["MLP +Dropout+AdamOptimizer(with different size of hidden layers(600,200))|96.48", 96.48])

x.add_row(["MLP+Dropout+Adam Optimizer with 3 hidden layers", 97.48])
x.add_row(["MLP +Dropout+Adam Optimizer with 3 hidden layers", 96.09])

print(x)

```

Model	Accuracy
Softmax classifier	90.02
MLP+Sigmod activation+SGD Optimizer	88.05
MLP+Sigmoid activation+ADAM	97.95
MLP+ReLU+SGD	98.46
MLP+ReLU+ADAM	98.43
MLP+Batch Norm on hidden Layers+Adam Optimizer	97.37
MLP+Dropout+AdamOptimizer	96.92
MLP +Dropout + AdamOptimizer(different size of hidden layers(540,360))	97.05
MLP +Dropout+AdamOptimizer(with different size of hidden layers(600,200)) 96.48	96.48
MLP+Dropout+Adam Optimizer with 3 hidden layers	97.48
MLP +Dropout+Adam Optimizer with 3 hidden layers	96.09

Conclusion

1. On the mnist dataset sigmoid alongwith SGD optimiser is used which does not give much accuracy
2. On changing the optimiser keeping the activation unit same accuracy has increased
3. Again on changing activation unit to Relu and using SGD as an optimiser accuracy has improved
4. The accuracy decreased bit on replacing optimiser with adam
5. On using dropout the accuracy did not increase much, may be because dataset is small
6. On using dropout along with 2 and 3 layers accuracy has improved.